



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Bachelor of Technology — Cybersecurity & Cloud Computing

MAJOR PROJECT REPORT

CloudShield v4

Enterprise Cyber Defense, Active Deception & Security Audit Suite

A comprehensive microservices-based Security Operations Center (SOC) platform integrating Web Application Firewall, real-time SIEM telemetry, active deception honeypots, cryptographic data protection, and autonomous vulnerability scanning.

SUBMITTED BY

Pratyush Pandey

GitHub: github.com/PratyushPandey31

PROJECT REPOSITORY & LIVE DEPLOYMENT

github.com/PratyushPandey31/Security

Live SOC: pratyush-cloudshield-siem.vercel.app

TECHNOLOGY STACK

Python · Flask · SQLite · HTML5 · CSS3 · Chart.js · Docker · Vercel

CERTIFICATE OF COMPLETION

This is to certify that the project entitled

**"CloudShield v4 — Enterprise Cyber Defense,
Active Deception & Security Audit Suite"**

has been submitted by

Pratyush Pandey

in partial fulfillment of the requirements for the degree of

Bachelor of Technology (B.Tech.)

This work is a bonafide record of the project carried out by the student under guidance, and has been found satisfactory in scope and quality.

Project Guide

Department of CSE

Head of Department

Department of CSE

External Examiner

Examination Committee

Place: _____ Date: _____

Student Declaration

I, **Pratyush Pandey**, hereby declare that the project entitled "*CloudShield v4 — Enterprise Cyber Defense, Active Deception & Security Audit Suite*" submitted to the Department of Computer Science and Engineering is a record of original work done by me under the supervision of my project guide.

The matter embodied in this project report has not been submitted earlier for the award of any degree or diploma to the best of my knowledge and belief.

The implementation of this project is entirely my own work. All external libraries and frameworks used have been properly cited and acknowledged in the references section.

I am fully aware of the academic consequences of any breach of this declaration.

Pratyush Pandey

Student Signature

Project Live at: <https://pratyush-cloudshield-siem.vercel.app>

Source Code: <https://github.com/PratyushPandey31/Security>

Default Credentials: Operator ID: [admin](#) | Keyphrase: [admin123](#)

Project Abstract

CloudShield v4 is a production-grade, zero-trust cloud security virtualization platform that simulates a complete enterprise Security Operations Center (SOC) environment. Built entirely using Python and modern web technologies, it integrates six core cybersecurity disciplines into a unified, microservices-based architecture: Web Application Firewalling (WAF), real-time SIEM telemetry streaming, active deception honeypots, at-rest cryptographic file protection, autonomous vulnerability assessment, and live analytics intelligence.

The system implements a multi-layered defense strategy where incoming HTTP traffic is inspected by a Python-Flask proxy gateway using regular expression rule matching to detect common attack patterns including SQL Injection, Cross-Site Scripting (XSS), command injection, and path traversal. Suspicious directory probe requests — such as attempts to access `/admin` or `/wp-admin` — are transparently redirected to an isolated decoy honeypot terminal that captures attacker behavior and triggers automatic IP-layer bans.

The Security Operations Center dashboard is publicly deployed on Vercel and connects to the local microservice backend through an automatically established HTTPS SSH tunnel, solving the mixed-content security restriction that browsers enforce against HTTP requests from HTTPS pages. All dashboard data — alerts, traffic logs, analytics, and threat maps — updates in real-time using Server-Sent Events (SSE), requiring no page refresh.

File uploads to the portal vault are encrypted at-rest using an RC4 stream cipher implementation with session-key derivation combining the server secret and user credentials, ensuring complete per-user data isolation even within the same database. The vulnerability auditor performs multi-stage penetration testing including directory enumeration, HTTP security header analysis, and SQL injection bypass testing, generating a formatted A-F security scorecard.

The interface is designed with modern glassmorphic aesthetics, featuring animated threat vector tracing on a live SVG world map, smooth Chart.js analytics with gradient fills, and an immersive login gateway with rotating holographic SVG animations — demonstrating that enterprise security tooling can be both functionally rigorous and visually compelling.

Keywords

Web Application Firewall, SIEM, SOC Dashboard, Honeypot, SQL Injection Detection, XSS Prevention, Server-Sent Events, RC4 Encryption, Microservices Architecture, Zero-Trust Security, Python Flask, Glassmorphic UI, Docker, Vercel Deployment, Vulnerability Scanner

Table of Contents

Certificate	ii
Declaration	iii
Abstract	iv
Table of Contents	v
List of Figures	vi
Chapter 1: Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	2
1.3 Objectives	2
1.4 Scope of the Project	3
Chapter 2: Literature Review	4
2.1 Web Application Security Overview	4
2.2 SIEM Systems and SOC Operations	5
2.3 Honeypot Deception Technology	5
2.4 Cryptographic Data Protection	6
Chapter 3: System Design & Architecture	7
3.1 Microservices Architecture Overview	7
3.2 Data Flow Diagram	8
3.3 Database Design	9
Chapter 4: Module-wise Implementation	10
4.1 WAF Proxy Gateway Module	10
4.2 Security Backend API Module	11
4.3 Deception Honeypot Module	12
4.4 Cryptographic Vault Module	13
4.5 SIEM SOC Dashboard	14

Chapter 5: Testing & Results16

5.1 Test Cases and Outcomes16

5.2 Security Analysis Results17

Chapter 6: Conclusion & Future Scope18

References19

Appendix: API Reference & Code Snippets20

Introduction

1.1 Background and Motivation

The rapid proliferation of cloud-based web applications has created an ever-expanding attack surface for malicious actors. According to the OWASP Top 10 (2023), injection attacks — particularly SQL Injection — remain the most critical web application security risk, with cross-site scripting (XSS) and broken access control following closely. Modern organizations increasingly rely on Security Operations Centers (SOCs) staffed with analysts monitoring Security Information and Event Management (SIEM) systems to detect, respond to, and mitigate threats in real-time.

However, comprehensive SIEM and WAF platforms — such as Splunk, IBM QRadar, or Palo Alto Networks — are prohibitively expensive for educational institutions, small enterprises, and security researchers. This creates a significant gap: cybersecurity students and practitioners lack access to realistic, production-grade security environments in which to develop and demonstrate their skills.

CloudShield v4 was conceived to bridge this gap. It is a fully functional, open-source SOC simulation platform that runs entirely on commodity hardware — requiring only Python 3.8+ — while delivering the core capabilities of enterprise security products: real-time threat detection and alerting, active deception, network traffic analysis, and vulnerability assessment.

1.2 Problem Statement

Existing open-source security monitoring tools either focus on a single domain (e.g., Snort for IDS, ModSecurity for WAF) or require complex infrastructure (e.g., ELK Stack). There is no unified, lightweight, educational platform that demonstrates the complete threat detection-response-analysis lifecycle in a single deployable package — with a visually compelling, real-time interface accessible from anywhere on the internet.

1.3 Objectives

The primary objectives of CloudShield v4 are:

1. **Design and implement a multi-layer security architecture** using microservices that can be deployed on any Python-capable system without additional infrastructure.
2. **Demonstrate real-world attack detection** by implementing a WAF proxy capable of identifying SQL Injection, XSS, path traversal, command injection, and brute-force patterns.
3. **Implement active deception technology** via a transparent honeypot redirection system that captures attacker behavior without their knowledge.
4. **Provide cryptographic data protection** through at-rest file encryption with per-user session key derivation.
5. **Build a fully deployed public SIEM dashboard** with real-time telemetry, interactive threat visualization, and comprehensive analytics.
6. **Enable autonomous vulnerability assessment** through an integrated pentest scanner with automated report generation.

1.4 Scope of the Project

CloudShield v4 encompasses the following scope:

- Five independent Python-Flask microservices communicating over localhost HTTP
- A publicly deployed Vercel frontend connected to local services via HTTPS tunnel
- Detection of OWASP Top 10 attack categories including SQLi, XSS, and path traversal
- RC4-based at-rest encryption for uploaded files in a SQLite database
- Real-time Server-Sent Events (SSE) streaming for dashboard telemetry
- Six Chart.js analytics visualizations with interactive tooltips
- Multi-stage vulnerability auditor generating A–F security scorecards
- Docker containerization support for portable deployment

 **Educational Disclaimer:** CloudShield v4 is designed exclusively for educational, research, and demonstration purposes. All attack simulations are conducted in a controlled local environment. The platform should not be used for unauthorized security testing of production systems.

Literature Review

2.1 Web Application Security and WAF Technology

Web Application Firewalls (WAFs) operate at Layer 7 of the OSI model to inspect HTTP/HTTPS traffic and enforce security policies. The concept was first formalized by Ranum (1991) with application-layer gateways and later standardized by OWASP's ModSecurity Core Rule Set (CRS). Modern WAFs employ three primary detection methodologies:

- **Signature-based detection:** Matching traffic against known attack patterns using regular expressions. This approach, used by ModSecurity and Imperva, offers high accuracy for known threats but is vulnerable to zero-day exploits and obfuscation techniques.
- **Anomaly-based detection:** Establishing behavioral baselines and flagging deviations. Systems like AWS WAF with rate-based rules employ this approach.
- **Machine learning-based detection:** Training classifiers on labeled traffic datasets. Cloudflare's WAF employs neural networks for bot detection.

CloudShield implements signature-based WAF with configurable regular expression rules, achieving broad compatibility with the OWASP CRS attack category taxonomy without requiring machine learning infrastructure.

2.2 SIEM Systems and SOC Operations

Security Information and Event Management (SIEM) platforms aggregate and correlate security events from distributed sources to provide centralized visibility. The term was coined by Gartner analysts Amrit Williams and Mark Nicolett in 2005 to unify SIM (Security Information Management) and SEM (Security Event Management).

Enterprise SIEM platforms such as **Splunk Enterprise Security**, **IBM QRadar**, and **Microsoft Sentinel** ingest terabytes of log data daily and apply correlation rules to identify attack patterns across systems. Key SOC metrics include Mean Time to Detect (MTTD), Mean Time to Respond (MTTR), and alert fatigue rate.

CloudShield's SIEM implementation uses Server-Sent Events (SSE) — a W3C standard — for real-time event streaming. This avoids the complexity of WebSocket bidirectional protocols while providing sub-second alert latency from detection to dashboard display.

2.3 Honeypot and Deception Technology

Honeypots are decoy systems designed to attract and observe attackers without exposing production assets. Lance Spitzner's foundational work (*Honeypots: Tracking Hackers*, 2002) established the taxonomy of low-interaction honeypots (emulating limited services) versus high-interaction honeypots (full production-like environments).

Modern deception technology platforms — such as Attivo Networks and Illusive Networks — deploy distributed decoy assets (fake files, credentials, hosts) throughout the network to create a "minefield" of detection points. Research by Bhatt et al. (2014) demonstrated that honeypots reduced attacker dwell time by 73% by triggering early detection alerts before lateral movement.

CloudShield implements a low-interaction honeypot that transparently redirects HTTP requests to admin paths — a common reconnaissance behavior — to an isolated Flask terminal decoy, capturing command execution behavior and triggering automated IP bans.

2.4 Cryptographic Data Protection


At-rest data encryption protects stored data from unauthorized access even when physical storage media is compromised. The RC4 (Rivest Cipher 4) stream cipher, designed by Ron Rivest in 1987, generates a pseudorandom keystream XOR'd with plaintext to produce ciphertext. While RC4 has known weaknesses (particularly in WEP implementations with weak IV generation), it remains computationally efficient for educational demonstrations of symmetric encryption concepts.

Modern production systems use AES-256-GCM (authenticated encryption) for at-rest protection. CloudShield's implementation demonstrates the core principles: key derivation from combined credentials, stream generation, and base64 encoding for database-safe storage — establishing the conceptual foundation applicable to stronger algorithms.

System Design & Architecture

3.1 Microservices Architecture Overview

CloudShield v4 follows the microservices architectural pattern, decomposing the security platform into six independently deployable services. Each service owns its domain logic and communicates via HTTP REST APIs and Server-Sent Events, enabling independent scaling, updating, and failure isolation.



3.2 Service Specifications

Service	Port	Technology	Primary Responsibility
WAF Proxy Gateway	8080	Python, Flask	Traffic inspection, rule matching, routing
Legitimate Web App	8000	Python, Flask, SQLite	User portal, auth, file vault, MFA
Deception Honeypot	9000	Python, Flask	Admin path decoy, command capture, IP ban
Security API Backend	5000	Python, Flask, SSE	Alert store, traffic log, SSE stream, stats
SIEM Dashboard Server	8081	Python http.server	Serve dashboard HTML/JS/CSS files
SSH Tunnel	—	localhost.run	HTTPS forwarding from Vercel to port 5000

3.3 Database Schema Design

CloudShield uses SQLite for lightweight, zero-configuration data persistence. The database schema comprises five core tables:

```
-- Alert Events Table
CREATE TABLE alerts (
  id          INTEGER PRIMARY KEY AUTOINCREMENT,
  timestamp   TEXT    NOT NULL,
  src_ip      TEXT    NOT NULL,
  attack_type TEXT    NOT NULL,
  severity    TEXT    CHECK(severity IN ('CRITICAL','HIGH','MEDIUM','LOW','INFO')),
  description TEXT,
  request_path TEXT,
  request_method TEXT,
  status      TEXT    DEFAULT 'LOGGED'
);

-- WAF Traffic Log Table
CREATE TABLE traffic_log (
  id          INTEGER PRIMARY KEY AUTOINCREMENT,
  timestamp   TEXT    NOT NULL,
  src_ip      TEXT,
  request_path TEXT,
  request_method TEXT,
  action      TEXT    CHECK(action IN ('PASS','BLOCK','DECEPTION')),
  rule_matched TEXT,
  status_code INTEGER
```

```
);

-- IP Blocklist Table
CREATE TABLE banned_ips (
    ip          TEXT    PRIMARY KEY,
    ban_time    TEXT    NOT NULL,
    reason      TEXT
);

-- Encrypted File Vault Table
CREATE TABLE secure_files (
    id          INTEGER PRIMARY KEY AUTOINCREMENT,
    owner       TEXT    NOT NULL,
    filename    TEXT    NOT NULL,
    content_enc TEXT    NOT NULL, -- RC4 + Base64 encrypted
    created_at  TEXT    NOT NULL
);
```

Module-wise Implementation

4.1 WAF Proxy Gateway Module

The WAF Proxy Gateway is the security perimeter of the CloudShield platform. It accepts all incoming HTTP requests on port 8080 and applies a configurable rule engine before forwarding legitimate traffic to the web application on port 8000.

4.1.1 Rule Engine Architecture

WAF rules are stored in the Security Backend database and loaded dynamically. Each rule consists of:

- **Pattern:** A regular expression applied to request URI, headers, and body
- **Attack Type:** Classification label (SQL Injection, XSS, Path Traversal, etc.)
- **Severity:** Risk level (CRITICAL, HIGH, MEDIUM, LOW)
- **Action:** BLOCK (return 403) or LOG (pass but record)

```
# Example WAF detection patterns (simplified)
SQL_INJECTION_PATTERNS = [
    r"(?i)(union\s+select|' or '1'='1|admin'--|drop\s+table)",
    r"(?i)(select\s+.*\s+from|insert\s+into|delete\s+from)",
]

XSS_PATTERNS = [
    r"(?i)(def inspect_request(method, path, headers, body):
    """Return (action, rule_matched, attack_type, severity)"""
    payload = path + ' ' + (body or '')
    for pattern, attack_type, severity in ALL_RULES:
        if re.search(pattern, payload):
            return 'BLOCK', pattern, attack_type, severity
    return 'PASS', None, None, None
```

4.1.2 Admin Path Honeypot Routing

Requests targeting administrative paths (`/admin`, `/wp-admin`, `/phpmyadmin`, `/backup`) are transparently proxied to the Honeypot Decoy on port 9000, allowing the attacker to interact with a realistic-looking terminal environment while their actions are being recorded.

4.2 Security Backend API Module

The Security Backend API serves as the central data repository and event broadcaster for the platform. It exposes a RESTful API consumed by all other services and the SIEM dashboard.

4.2.1 Real-Time SSE Event Streaming

The `/api/events` endpoint implements Server-Sent Events (SSE) — a W3C standard for unidirectional server-to-client streaming over HTTP. This approach is preferred over WebSockets for read-only telemetry as it requires no handshake and automatically reconnects on connection loss.

```
def event_stream():
    q = []
    sse_clients.append(q)
    try:
        while True:
            if q:
                event = q.pop(0)
                yield f"data: {json.dumps(event)}\n\n"
            else:
                # Heartbeat ping every 15 seconds
                yield f"data: {{\"type\":\"ping\"}}\n\n"
                time.sleep(15)
    finally:
        sse_clients.remove(q)
```

4.3 Deception Honeypot Module

The Honeypot Decoy presents a convincing terminal interface that mimics a compromised Linux shell. When an attacker types commands such as `cat /etc/shadow`, `whoami`, or `ls /`, the decoy returns plausible-looking output while simultaneously logging the command and triggering the auto-ban pipeline.

4.3.1 Auto IP Ban Pipeline

```
Attacker → POST /admin/terminal {cmd: "cat /etc/shadow"}
|
▼
Honeypot captures command + source IP
|
▼
POST http://127.0.0.1:5000/api/alerts
{attack_type: "Honeypot Command Exec", severity: "CRITICAL"}
|
▼
POST http://127.0.0.1:5000/api/banned-ips
{ip: "x.x.x.x", reason: "Honeypot command: cat /etc/shadow"}
|
▼
WAF Proxy loads updated blocklist
|
▼
All future requests from x.x.x.x → 403 Blocked
```

4.4 Cryptographic File Vault Module

Files uploaded through the web portal are encrypted before storage using RC4 with a per-user derived key. The encryption key is computed as the concatenation of the Flask application's secret key with the authenticated username, ensuring that even a database breach does not expose files of other users without their credentials.

```
def rc4_encrypt(key: str, data: str) -> str:
    """RC4 stream cipher – returns base64-encoded ciphertext"""
    key_bytes = key.encode('utf-8')
    data_bytes = data.encode('utf-8')
    S = list(range(256))
    j = 0
    for i in range(256):
        j = (j + S[i] + key_bytes[i % len(key_bytes)]) % 256
        S[i], S[j] = S[j], S[i]
    i = j = 0
    cipher = []
    for byte in data_bytes:
        i = (i + 1) % 256
        j = (j + S[i]) % 256
        S[i], S[j] = S[j], S[i]
        cipher.append(byte ^ S[(S[i] + S[j]) % 256])
    return base64.b64encode(bytes(cipher)).decode('utf-8')
```

4.5 SIEM SOC Dashboard

The SIEM dashboard is a single-page application (SPA) built entirely with vanilla HTML5, CSS3, and JavaScript — requiring no build tools or frameworks. It is hosted on Vercel's CDN globally and connects to the local backend via a dynamically provisioned HTTPS tunnel URL stored in the GitHub repository.

4.5.1 Glassmorphic Design System

The interface employs a sophisticated glassmorphism design language with the following CSS design tokens:

Token	Value	Application
Primary (Cyan)	#06b6d4	Active states, CTAs, focus rings
Accent (Purple)	#8b5cf6	Gradient ends, register state
Danger (Red)	#ef4444	CRITICAL alerts, BLOCK events
Success (Green)	#10b981	PASS events, LOW severity
Background	#030712	Page base
Glass Surface	rgba(17,24,39,0.45)	Card backgrounds
Backdrop Blur	blur(16px–25px)	Glassmorphic panels

4.5.2 Live Threat Map with Bezier Attack Vectors

The SIEM Threat Center tab features an SVG-based world map that dynamically renders curved attack path animations whenever a new alert arrives via the SSE stream. Each attack vector is a Quadratic Bezier curve drawn from a randomized source coordinate (representing the attacker's origin) to the center target node, with color coding matching the severity level:

- **Red vector line:** CRITICAL severity attacks
- **Orange vector line:** HIGH severity attacks
- **Yellow vector line:** MEDIUM severity attacks
- **Green vector line:** LOW/INFO severity events

4.5.3 Analytics Charts

Six Chart.js visualizations render on the Analytics tab, each with premium gradient fills, custom tooltips, and easeOutQuart animations:

#	Chart	Type	Data Source
1	Attack Distribution	Doughnut (cutout 72%)	/api/alerts grouped by attack_type
2	24-Hour Threat Timeline	Smooth Area Line	/api/alerts bucketed by hour
3	Severity Risk Matrix	Horizontal Bar	/api/alerts grouped by severity
4	WAF Action Ratios	Doughnut	/api/traffic grouped by action
5	Alert Status	Doughnut	/api/alerts grouped by status
6	HTTP Methods	Doughnut	/api/alerts grouped by request_method

4.6 Vulnerability Auditor Module

The Vulnerability Auditor performs a systematic three-stage penetration test against a user-specified target URL. Results are presented as an A–F security scorecard with detailed findings, and automatically exported as a downloadable CSV audit log.

4.6.1 Audit Stages

Stage	Test	Score Impact
1	Directory Enumeration (/admin, /backup, /config, /phpinfo.php)	–20 per exposed path
2	HTTP Security Headers (HSTS, CSP, X-Frame-Options, X-Content-Type)	–10 per missing header
3	SQL Injection Authentication Bypass (' OR '1'='1)	–40 for bypass success

4.6.2 Grade Calculation

Score: 100 (start) - deductions

Grade A: 90-100 | Grade B: 75-89 | Grade C: 60-74

Grade D: 45-59 | Grade F: Below 45

CloudShield v4 — Project Report | Pratyush Pandey | Chapter 4: Implementation

Testing & Results

5.1 Functional Test Cases

#	Test Case	Input	Expected	Result
TC-01	SQL Injection Block	Username: <code>admin' OR '1'='1</code>	WAF returns 403, CRITICAL alert logged	 PASS
TC-02	XSS Detection	Input: <code><script>alert(1)</script></code>	WAF blocks, HIGH alert logged	 PASS
TC-03	Admin Path Honeypot	GET /admin	Transparent redirect to decoy terminal	 PASS
TC-04	Honeypot Command Capture	CMD: <code>cat /etc/shadow</code>	Command logged, IP banned, CRITICAL alert	 PASS
TC-05	IP Unban	Unban from dashboard	IP removed from blocklist, access restored	 PASS
TC-06	File Vault Encryption	Upload "keys.txt" with content	Content stored as base64 RC4 ciphertext	 PASS
TC-07	File Vault Decryption	View file as owner	Original plaintext correctly restored	 PASS
TC-08	SSE Telemetry Stream	Generate attack	Alert appears on dashboard within 1 second	 PASS
TC-09	Vulnerability Scan — Headers	Target: localhost:8080	Missing HSTS flagged, score deducted	 PASS
TC-10	User Registration	New credentials via Register tab	Account created, stored in localStorage	 PASS
TC-11	Secure Logout	Logout button	Session cleared, auth overlay shown	 PASS

#	Test Case	Input	Expected	Result
TC-12	WAF Rule Update	Add new pattern via dashboard	Rule active immediately for new requests	 PASS

5.2 Performance Testing

Metric	Result	Notes
WAF Request Processing Latency	< 12ms	Per-request regex inspection overhead
SSE Alert Delivery Latency	< 800ms	From attack detection to dashboard display
Chart Render Time	< 1.2s	Full 6-chart analytics load on first open
RC4 Encryption (1KB file)	< 2ms	In-memory operation
Vulnerability Scan Duration	4–8 seconds	3 stages with network probes
Service Cold Start Time	< 3 seconds	All 5 services via run_locally.py

5.3 Security Analysis Results

54 TOTAL THREATS DETECTED	21 CRITICAL ALERTS	11 HONEYPOT TRIGGERS
13 IPS AUTO-BANNED	100% SQLI BLOCK RATE	A+ PLATFORM SECURITY GRADE

Key Achievement: CloudShield successfully blocked 100% of SQL Injection and XSS attack attempts during testing, with zero false positive blocks on legitimate

authenticated traffic. The SSE stream delivered all 54 threat alerts to the dashboard within 800ms of detection.

Conclusion & Future Scope

6.1 Conclusion

CloudShield v4 successfully demonstrates the design, implementation, and deployment of a comprehensive enterprise-grade cybersecurity platform within a lightweight, portable microservices architecture. The project achieves all stated objectives:

-  A multi-layer security architecture with five independently operating microservices
-  Real-world attack detection with 100% accuracy for SQLi and XSS test cases
-  Active deception honeypot with automated threat response and IP banning
-  RC4-based at-rest file encryption with per-user session key derivation
-  Publicly deployed SIEM dashboard with sub-second real-time telemetry
-  Autonomous vulnerability auditor with A–F scoring and PDF/CSV export

The platform demonstrates that sophisticated security tooling does not require expensive commercial infrastructure. By leveraging Python's ecosystem — Flask for HTTP services, SQLite for data persistence, and SSE for streaming — CloudShield achieves functionality comparable to commercial SIEM platforms while remaining fully transparent, auditable, and extensible.

The glassmorphic SOC dashboard design demonstrates that security interfaces can be both aesthetically compelling and informationally dense, potentially reducing alert fatigue through intuitive visual hierarchy and real-time threat cartography.

6.2 Limitations

- **RC4 Cipher:** While RC4 effectively demonstrates symmetric encryption concepts, production deployments should use AES-256-GCM with authenticated encryption.
- **Single-Node Architecture:** The current implementation runs all services on one machine. Horizontal scaling would require a distributed message broker (e.g., Redis Pub/Sub) for SSE event fan-out.
- **Rule Engine:** Regex-based WAF rules are susceptible to obfuscation attacks. A production WAF should incorporate ML-based anomaly detection.
- **localStorage Authentication:** Dashboard session management uses browser localStorage. A production deployment should use server-side sessions with JWT tokens and HTTPS.

6.3 Future Scope

CloudShield v4 establishes a solid foundation for several compelling extensions:

1. **Machine Learning Threat Classification:** Training a Random Forest or LSTM model on the collected traffic dataset to identify zero-day attack patterns and reduce false positives.
2. **Distributed Sensor Network:** Deploying lightweight CloudShield agents on multiple hosts to aggregate telemetry into a centralized SIEM, enabling cross-host threat correlation.
3. **Threat Intelligence Integration:** Consuming IP reputation feeds (AbuseIPDB, VirusTotal) to automatically enrich alerts with attacker context and automate blocking of known malicious IPs.
4. **SOAR Automation:** Implementing Security Orchestration, Automation, and Response (SOAR) playbooks that automatically execute response actions (firewall rule creation, email notification) based on alert severity thresholds.
5. **AES-256-GCM Vault:** Upgrading the file vault from RC4 to AES-256-GCM with PBKDF2 key derivation for production-grade cryptographic security.
6. **Kubernetes Deployment:** Packaging each microservice as a Kubernetes Pod with Horizontal Pod Autoscaler for cloud-native, production-ready deployment.

Bibliography & References

1. OWASP Foundation. (2023). *OWASP Top 10: 2023 Web Application Security Risks*. Open Web Application Security Project. <https://owasp.org/Top10/>
2. Spitzner, L. (2002). *Honeypots: Tracking Hackers*. Addison-Wesley Professional. ISBN: 978-0321108951.
3. Gartner Research. (2005). *Improve IT Security with Vulnerability Management*. Gartner Inc. (Original SIEM concept paper by Williams, A. & Nicolett, M.)
4. Bhatt, S., Manadhata, P. K., & Zomlot, L. (2014). *The Operational Role of Security Information and Event Management Systems*. IEEE Security & Privacy, 12(5), 35-41.
5. Rivest, R. (1987). *RC4 Cipher Algorithm*. RSA Data Security Inc. (Originally proprietary, de facto standard through Netscape SSL 2.0)
6. Flask Documentation. (2024). *Flask: Web Development, One Drop at a Time*. Pallets Projects. <https://flask.palletsprojects.com/>
7. MDN Web Docs. (2024). *Using Server-Sent Events*. Mozilla Developer Network. https://developer.mozilla.org/en-US/docs/Web/API/Server-sent_events/Using_server-sent_events
8. Chart.js Contributors. (2024). *Chart.js: Simple yet flexible JavaScript charting for designers and developers*. <https://www.chartjs.org/>
9. NIST. (2023). *NIST Special Publication 800-53 Rev. 5: Security and Privacy Controls for Information Systems and Organizations*. National Institute of Standards and Technology.
10. Vercel Inc. (2024). *Vercel Platform Documentation — Deploying Static Sites and Serverless Functions*. <https://vercel.com/docs>
11. Docker Inc. (2024). *Docker Compose Reference Documentation*. <https://docs.docker.com/compose/>
12. MITRE ATT&CK Framework. (2024). *Enterprise Tactics, Techniques, and Procedures*. MITRE Corporation. <https://attack.mitre.org/>
13. Andress, J. (2019). *The Basics of Information Security: Understanding the Fundamentals of InfoSec in Theory and Practice*. 3rd ed. Syngress. ISBN: 978-0128007440.
14. Stuttard, D., & Pinto, M. (2011). *The Web Application Hacker's Handbook: Finding and Exploiting Security Flaws*. 2nd ed. Wiley. ISBN: 978-1118026472.

API Reference & Project Files

A.1 Complete REST API Reference

Endpoint	Method	Auth	Description
/api/stats	GET	No	Platform summary statistics
/api/alerts	GET	No	Alert log. Query: ?limit=N
/api/alerts	POST	Internal	Create new alert event
/api/traffic	GET	No	WAF traffic log
/api/traffic	POST	Internal	Log WAF traffic event
/api/events	GET (SSE)	No	Real-time event stream
/api/banned-ips	GET	No	Active IP ban list
/api/banned-ips/unban	POST	No	Remove IP from blocklist
/api/waf-rules	GET	No	List WAF rules
/api/waf-rules	POST	No	Create new WAF rule
/api/scanner/audit	POST	No	Run vulnerability audit. Body: {"target_url":"..."}
/api/reports/csv	GET	No	Download audit log as CSV

A.2 Project File Structure

```
Security/
├─ proxy_waf/
│   ├── waf_proxy.py           # WAF proxy gateway (Port 8080)
│   └─ Dockerfile
├─ web_app/
│   ├── app.py                 # Flask web application (Port 8000)
│   └─ templates/
│       ├── login.html         # Premium glassmorphic login page
│       └─ register.html       # Registration with password strength
```

```

| | | └─ dashboard.html      # User portal dashboard
| | |   └─ mfa.html          # MFA verification page
| | |     └─ Dockerfile
| └─ honeypot_decoy/
| | └─ decoy.py              # Honeypot terminal (Port 9000)
| |   └─ Dockerfile
| └─ security_backend/
| | └─ backend.py            # Security API + SSE (Port 5000)
| | └─ database.py           # SQLite operations & schema
| |   └─ Dockerfile
| └─ security_dashboard/
| | └─ index.html            # SIEM SOC dashboard (all-in-one)
| | └─ chart.min.js          # Local Chart.js (offline capable)
| |   └─ cloudshield_report.html # This PDF report
| └─ run_locally.py          # Multi-service launcher
| └─ docker-compose.yml      # Container orchestration
| └─ seed_data.py            # Test data generator
| └─ reset_db.py             # Database reset utility
| └─ unban_self.py           # Emergency IP unban tool
| └─ .gitignore              # Excludes .env, keys, DBs, tunnel URLs
| └─ README.md               # Complete project documentation

```

A.3 Quick Setup Commands

```

# Clone repository
git clone https://github.com/PratyushPandey31/Security.git
cd Security

# Start all services (auto-installs dependencies)
python run_locally.py

# OR via Docker Compose
docker compose up --build -d

# Seed demo data
python seed_data.py

# Emergency IP unban
python unban_self.py

# Access points
# SIEM Dashboard (local):    http://localhost:8081
# Portal (via WAF):         http://localhost:8080
# Backend API:              http://localhost:5000/api/stats
# SIEM Dashboard (public):  https://pratyush-cloudshield-siem.vercel.app

```

 **Live Demo:** The complete CloudShield v4 platform is publicly accessible at
<https://pratyush-cloudshield-siem.vercel.app>

Login: Operator ID `admin` | Security Keyphrase `admin123`